

Overview

- Notebook builds a lightweight, fully analytic demo of a "liquid" neuron (single-step LTC) trained to invert noise on synthetic expectation values, entirely with NumPy and manual gradients, no quantum runtime or autograd involved (notebooks/demo_math.ipynb:1).

Data Pipeline

- Synthetic targets come from drawing random angles θ in $[0, 2\pi)$ and mapping to ideal expectations $y = \cos(\theta)$ (notebooks/demo_math.ipynb:38-45).
- Noisy observations u are produced via an affine shrink + bias plus Gaussian shot noise, then clipped back to the physical interval $[-1, 1]$, mimicking amplitude damping toward +1 (notebooks/demo_math.ipynb:46-55).
- Dataset is shuffled and split 80/20 into train/val splits (notebooks/demo_math.ipynb:56-60).

Model Structure

- The forward routine implements a single LTC update: starting from $x_0 = u$, it computes the gating term $g = \tanh(wr \cdot x_0 + w \cdot u + b)$, evolves one Euler step $x_1 = (x_0 + dt \cdot g \cdot A) / (1 + dt \cdot (1/\tau + g))$, and projects the readout through $\tanh(v \cdot x_1 + c)$ to stay within observable bounds (notebooks/demo_math.ipynb:62-82).
- Helper mse returns batch mean-squared error (notebooks/demo_math.ipynb:81).

Manual Gradients

- gradients caches activations from forward, then backpropagates analytically:
 - Starts with $dL/dy_{\text{hat}} = 2 \cdot (y_{\text{hat}} - y) / N$ for the batch (notebooks/demo_math.ipynb:84-91).
 - Uses $\tanh' = 1 - \tanh^2$ to differentiate the output stage for parameters v and c (notebooks/demo_math.ipynb:92-101).
 - Applies quotient-rule sensitivities for $x_1 = \text{num/den}$ with respect to g , A , and τ , including the τ -only term in the denominator (notebooks/demo_math.ipynb:102-124).
 - Chains g -derivatives back to weights w , wr , b while aggregating over the batch (notebooks/demo_math.ipynb:125-135).
 - Returns the 7-parameter gradient vector alongside the latest predictions (notebooks/demo_math.ipynb:134-136).

Training Loop

- Parameters are initialized near zero with $\tau > 0$, learning rate 0.02, 30 epochs, batch size 200, and updated via vanilla SGD (params -= lr * grads) with a post-update projection enforcing $\tau \geq 0.05$ (notebooks/demo_math.ipynb:138-160).
- After each epoch it logs train/val MSE using fresh forward passes, storing the curves for plotting (notebooks/demo_math.ipynb:161-166).
- A Matplotlib cell plots and saves training_loss_math.png, illustrating convergence of both splits (notebooks/demo_math.ipynb:167-178).

Validation Analysis

- A second cell compares mean absolute error on raw noisy validation data versus the mitigated predictions, printing both metrics and saving a bar chart (math_mae_comparison.png) that quantifies the improvement (notebooks/demo_math.ipynb:221-234).
- The final cell simply echoes y_{clean} , leaving the clean targets in the output for quick

inspection (notebooks/demo_math.ipynb:257).